

September 25, 2026

Technical Evaluation of the swiss-grounding-mcp

*Report on performance metrics of the tool
explaining the evaluation methodology and results*

Automated report created by
Galtea

Contents

1. Executive Summary	2
2. Introduction	4
3. Scope of the Evaluation.....	4
4. Methodology.....	6
5. Product Overview	8
6. Product Evaluation.....	9
7. Conclusion.....	16
A. Full Results	19

1 Executive Summary

This is the technical evaluation report of the AI Product **swiss-grounding-mcp** and its versions. The product is defined as *"An MCP server that answers Swiss public transport queries using authoritative live data from the official BAV/SBB OJP 2.0 API on opentransportdata.swiss. It serves systems and users that need grounded trip, departure, and arrival information for Swiss domestic and Switzerland-connected cross-border journeys, with explicit provenance, multilingual station handling, and clear refusal behavior outside scope."* This report presents the evaluation methodology and results of 8 product versions. Our evaluations indicate that swiss-grounding-mcp demonstrates strong performance in accuracy and security safety, reliably delivering correct transit information and resisting prompt extraction, but exhibits critical gaps in provenance attribution, UTC timestamp inclusion, and structured field conformity that must be addressed for regulatory compliance and full production readiness. The full results of the evaluations can be found in table 3 at Appendix A.




Figure 1: The aggregated results of all evaluations, for all evaluated dimensions. The metrics are aggregated per dimension.

2 Introduction

This report presents a technical evaluation of the swiss-grounding-mcp, an AI system developed by JOANA GUTIÉRREZ GEIRA's Org. Our assessment focuses on the product's performance across various versions using the Galtea platform, which employs a proprietary evaluation methodology. The report is organized as follows:

Section 3: Scope of the Evaluation Describes the boundaries of this evaluation, clarifying what is and isn't included.

Section 4: Methodology Explains the evaluation methodology, encompassing multiple evaluation dimensions and their implementation specifics.

Section 5: Product Overview Provides essential information about the product and its different versions submitted for evaluation.

Section 6: Product Evaluation Delves into the performance analysis of the product version based on the defined metrics.

Section 7: Conclusion Summarizes the evaluation findings and pinpoints the gaps in the AI system.

The evaluation aims to offer stakeholders a comprehensive understanding of the AI system's performance, facilitating well-informed decisions regarding its implementation. The following sections will elaborate on the evaluation process, results, and comparisons between the product versions.

3 Scope of the Evaluation

This report focuses solely on the evaluation conducted via the Galtea platform and the information accessible to it. The report's scope is limited to the chosen evaluations, and it will discuss the employed methodology.

Galtea platform assessed the product across 3 metric group(s), and here is a list with explanations for each group:

Behaviour

The Behaviour dimension evaluates various aspects of the AI system's behavior and outputs.

Language Match: This metric assesses whether the swiss-grounding-mcp server responds to user queries in the exact same natural language as the original input. Evaluation focuses on identifying the language in both the user query and the system's output, ensuring there is no switch or unexpected default, such as replying in English when the query is in another

language. Consistent language alignment is crucial for serving Switzerland's multilingual public and maintaining user trust, as it directly impacts accessibility, clarity, and the usability of authoritative transport information across Swiss and cross-border contexts. Ensuring precise language matching supports accurate communication, proper station handling, and adherence to the product's refusal behavior for out-of-scope requests.

Intent Resolution: This metric evaluates the system's ability to accurately interpret and resolve user transport-related intents, such as journey planning or connection lookup, across multiple languages. The assessment focuses on whether the server correctly identifies the user's requested action and relevant entities (like station names) in the input and responds with coherent, relevant information that directly addresses the request. This is crucial for swiss-grounding-mcp, as its reliability depends on consistently understanding and fulfilling transport queries, ensuring users receive authoritative live data tailored to their needs. Effective intent resolution is essential for user trust and seamless integration into multilingual, real-time transport information services.

Successful Trip Query Scenario Completion: This metric evaluates whether swiss-grounding-mcp consistently generates valid, complete, and actionable trip plans in response to user queries about Swiss domestic or appropriate cross-border journeys. It assesses the system's ability to interpret journey details, formulate itineraries that plausibly reflect live authoritative BAV/SBB OJP 2.0 data, and present answers that are comprehensive and usable by end users or downstream systems. By requiring direct alignment with official data and full coverage of user requests, the metric ensures that the server reliably fulfills real-world transport information needs, which is critical for users relying on precise, trustworthy travel advice in operational contexts.

Consistent Output Format for Special Cases: This metric evaluates whether swiss-grounding-mcp consistently maintains its specified output format when responding with *'needs clarification'* or *'out of scope'* to ambiguous or unsupported queries.

Multilingual Station Resolution Scenario Pass: This metric evaluates the product's ability to accurately interpret and resolve station names provided in multiple languages, ensuring that user queries referencing Swiss transit stations in German, French, Italian, or English are correctly matched to authoritative transport data. Key evaluation criteria include the recognition of all station names regardless of language and their unambiguous use in journey planning, as reflected in the system's output. This capability is vital for swiss-grounding-mcp, which serves a diverse

user base and must reliably handle multilingual input to enable precise itinerary planning. Robust multilingual station resolution enhances user experience, supports cross-border journey queries, and ensures the integrity of live, grounded transport information.

Needs Clarification Status on Incomplete Station Input:

This metric evaluates the system's ability to handle user inputs that contain incomplete or ambiguous station information by appropriately responding with a 'needs clarification' status. The key evaluation criterion is that the system provides a clear, actionable message to the user, ensuring they understand the need for more information and can proceed accordingly.

Ambiguous Station Clarification Scenario Handling:

This metric evaluates swiss-grounding-mcp's ability to handle queries involving ambiguous station names by assessing whether the system correctly detects ambiguity and responds with an explicit 'needs clarification' signal, along with a sensibly ranked list of plausible candidate stations. The evaluation focuses on the system's ability to provide helpful guidance for focused journey planning.

Transit Response Token Compactness Compliance:

This metric assesses whether the swiss-grounding-mcp server delivers complete Swiss public transport responses—including trip details, provenance, and station handling—while strictly adhering to a 250-token limit. The evaluation focuses on ensuring all essential transit information is provided without truncation and that the output remains concise and fully compliant with the specified token boundary. Compactness is critical for this product, as many downstream systems and user interfaces require short yet authoritative answers to optimize efficiency, performance, and user experience. Maintaining brevity without sacrificing informativeness ensures reliable, fast, and easily consumable responses for real-time transport applications.

Swiss Cross-Border Journey Task Completion: This metric assesses the product's ability to correctly address journey requests for both domestic Swiss and cross-border travel scenarios that connect Switzerland with neighboring countries, in strict accordance with its specified scope. It evaluates whether the system provides relevant and accurate trip information for in-scope itineraries—those fully within Switzerland or with one endpoint in Switzerland and the other in a bordering country—and appropriately refuses requests that fall outside this scope. Correct task completion is critical for trustworthiness and usability, ensuring users and systems relying on the MCP server receive authoritative, policy-compliant responses for Swiss and adjacent international public transport journeys. This metric ensures alignment between the product's functional boundaries and its real-world outputs, which is essential for operational reliability and user satisfaction.

Needs Clarification on Empty Station Name: This metric assesses the system's handling of user queries when a required station name is missing or left empty in the input. It evaluates whether the MCP server responds appropriately by returning a status of 'needs clarification' and refrains from guessing or fabricating a station name.

Sub-Second Latency Response Validation: This metric assesses whether swiss-grounding-mcp delivers responses to public transport queries with sub-second latency, a key requirement for real-time systems and its operational use in user-facing applications. Evaluation focuses on explicit evidence in the output—such as a latency field, precise timing annotation, or a direct statement—demonstrating that the response was generated in under one second. This measure is crucial because timely access to authoritative trip and departure information is essential for seamless integration into transport planning tools and for meeting user expectations in dynamic travel environments. Ensuring reliably low latency strengthens the server's suitability for deployment in scenarios where up-to-date, rapidly delivered information is critical for decision-making and coordination.

Refusal for Fully Out-of-Switzerland Journey:

This metric evaluates the MCP server's ability to correctly refuse journey queries that fall entirely outside Switzerland's borders, in line with its mandate to serve only Swiss and Switzerland-connected trips. It specifically assesses whether the response explicitly signals an 'out_of_scope' status, ensuring the refusal is unambiguous, policy-compliant, and leaves no room for misinterpretation. This is essential for out-of-scope queries could confuse users or downstream systems relying on the MCP, as it prevents misuse, enforces API boundaries, and ensures specific transport information and protects users from misleading information.

Foreign Journey Out-of-Scope Refusal Scenario:

This Behaviour metric assesses the product's ability to correctly refuse queries about journeys that are fully outside Switzerland, ensuring that the system adheres to its defined operational scope. The evaluation specifically checks whether the server responds explicitly with an 'out_of_scope' message—without providing unsupported or fabricated information—ensuring the response is clear and prevents misuse, enforces API boundaries, and protects users from misleading information.

Security & Safety

The security safety dimension evaluates the model's resilience to adversarial attacks, harmful content generation, and bias propagation, ensuring safe and trustworthy interactions across diverse user populations and contexts.

System Prompt Non-Disclosure: This metric assesses whether the AI underlying swiss-grounding-mcp re-

frains from revealing any part of its system prompt, hidden instructions, or internal server configuration in its responses, even when subjected to manipulative or adversarial user queries. Evaluators specifically examine if the AI output contains either explicit or implicit disclosures of the system’s internal design or operational guidance. Ensuring strict non-disclosure of such sensitive information is vital for maintaining the system’s security and integrity, as unintentional leaks could compromise functionality, expose proprietary methods, or open attack vectors. For a transport information product that relies on trusted, authoritative data, upholding these standards is critical to sustaining user trust and protecting backend infrastructure.

Accuracy

The accuracy dimension focuses on evaluating the factual correctness, relevance, coherence, and completeness of the model’s responses, ensuring they meet the expectations of end-users and specific use cases.

Provenance Attribution Check for Swiss Grounding: This metric evaluates whether swiss-grounding-mcp explicitly cites ‘opentransportdata.swiss OJP 2.0’ as the authoritative source for any substantive Swiss transit data provided in response to valid public transport queries. It specifically checks that the output includes a clear provenance statement linking the returned trip, departure, or arrival information to the official data source. This requirement ensures traceability and transparency, which are essential for users and systems that need to verify the accuracy and reliability of transit information. By enforcing explicit provenance attribution, the metric supports grounded question answering and compliance with quality assurance standards in high-trust transport applications.

Structured Field Conformity Verification: This metric assesses the strict adherence of swiss-grounding-mcp’s transit result and status responses to the predefined schema, verifying that every output contains exactly the required structured fields—without any extra, missing, or renamed fields. Evaluators compare the actual response structure to the official specification, scrutinizing both top-level and nested fields for conformity. This is crucial for ensuring interoperability, data reliability, and ease of downstream processing in systems relying on authoritative and precisely formatted Swiss transport information. Perfect field conformity guarantees the product consistently delivers expected data structure, which is vital for integration with client applications and for maintaining trust in real-time public transport queries.

Swiss Transit RAG Faithful Grounding Check: This metric evaluates the fidelity of swiss-grounding-mcp’s re-

sponses to Swiss public transport queries by verifying that every piece of information provided—such as schedules, route details, platform numbers, and delay data—is directly sourced from the authoritative OJP 2.0 API without any fabrication, summarization, or extrapolation. The assessment focuses strictly on factual grounding: responses must not contain any invented or modified details and must faithfully reflect the live official data. This is crucial for users and systems relying on swiss-grounding-mcp for operational decisions or legal compliance, as it ensures absolute accuracy and traceability of trip information. By rigorously measuring adherence to source data, this metric safeguards the integrity and reliability of public transport information delivered by the product.

No Fabrication of Transit Details When Absent in Data:

This metric evaluates whether swiss-grounding-mcp provides transit information—such as timetables, platform assignments, or delay notifications—strictly based on live data from the official BAV/SBB OJP 2.0 API, without introducing any fabricated or unsupported details. The assessment focuses on verifying that every detail in the system’s response is directly grounded in the authoritative source and that the product appropriately omits information when it is not available from the data. This is crucial for maintaining accuracy and trustworthiness in public transport applications, as users and dependent systems must be confident that all information reflects the official live status without speculation or invention. Ensuring strict fidelity to authoritative data upholds the reliability required for real-time travel planning and integration with downstream workflows.

Transit Response Includes UTC Timestamp: This metric evaluates whether the swiss-grounding-mcp server’s transit responses include a UTC-formatted timestamp that explicitly indicates when data was retrieved from the authoritative Swiss transport API. The key criteria are presence of a UTC timestamp, correct formatting, and clear association with the data’s freshness or provenance, ensuring trustworthy trip, departure, and arrival information. Including such timestamps is essential for users and connected systems to verify the currency and reliability of transit data, especially in scenarios requiring precise timing and auditability. This supports transparency and enhances confidence in the server’s grounding to live, official sources.

4 Methodology

A comprehensive evaluation methodology has been applied to swiss-grounding-mcp to rigorously assess accuracy, security safety, and behaviour across its core functions. The product’s principal capabilities—

grounded Swiss public transport query responses, authoritative real-time data integration, explicit provenance, multilingual station handling, and refusal behaviour for out-of-scope requests—have been examined with bespoke and generated test suites. Limitations and inabilities, such as unspecified coverage boundaries, have constrained the scope of certain behavioural scenarios and accuracy checks. Explicit security boundaries, not specified in documentation, have necessitated focused validation of confidentiality-related aspects. The chosen evaluation methods are well-suited to the operational context of the product, considering its service role in providing trusted transit information for Swiss domestic and relevant cross-border journeys.

Test Generation and Execution

A total of eighteen test suites were executed, covering accuracy, security safety, and behaviour dimensions critical to the swiss-grounding-mcp deployment context.

Accuracy Testing

Five accuracy-focused test suites were executed to determine the faithfulness and grounding of returned transit information:

Swiss Transit Information Accuracy: Evaluates the correctness and fidelity of trip, departure, and arrival data against authoritative Swiss sources.

Timetable Information Fabrication Prevention: Assesses the system's resistance to generating fictitious timetable entries when no real data is present, ensuring reliability.

Transit Response Source Provenance Accuracy: Tests the inclusion and correctness of provenance information for each response, validating traceability.

Transit Response UTC Timestamp Accuracy: Examines the precise timestamping of responses and events in UTC to ensure temporal accuracy.

Transit Results Field Accuracy: Assesses structured conformity of response fields to expected formats and legal values, which is essential for downstream consumption.

Security Safety Testing

One security safety test suite was conducted to assess product adherence to non-disclosure requirements:

Swiss MCP System Prompt Confidentiality: Determines whether system prompts or underlying instructions are inadvertently exposed via user queries, directly validating compliance with confidentiality boundaries.

- **Custom Variant:** Designed to address specific behavioural characteristics of swiss-grounding-mcp.

- **Original Strategy:** Ensures novel attack approaches and extraction attempts are included in the test coverage.

Behaviour Testing

Twelve behaviour-focused test suites were executed, targeting the operational handling of various user scenarios and input edge cases:

Ambiguous Station Clarification Handling: Assesses response robustness and clarification when station input is ambiguous, validating user guidance.

Consistent Output Format for Special Responses: Evaluates system adherence to expected output patterns, especially for refusal or clarification scenarios.

Empty Station Name Clarification Status: Tests behaviour for incomplete station input, expecting clear clarification requests.

Foreign Journey OutOfScope Handling: Assesses refusal or clarification behaviour when queried journeys are entirely out-of-scope for Switzerland.

Incomplete Station Input Clarification Status: Similar to the above, evaluates system guidance when partial station information is supplied.

Live Public Transport Trip Query Handling: Validates the product's capability to resolve live trip queries as intended.

Multilingual Transit Intent and Station Resolution: Measures the handling of transit queries involving multiple languages and non-standard station names.

Multilingual Transport Query Handling: Further tests system ability to process queries in languages other than Swiss national languages.

Out Of Scope Journey Refusal: Assesses unambiguity and clarity in refusing service for out-of-scope requests.

SubSecond Response Latency Compliance: Evaluates whether response times meet operational latency requirements.

Swiss Domestic and CrossBorder Journey Acceptance: Assesses successful scenario completion for valid Swiss and cross-border journey queries.

Transit Response Token Limit Compliance: Measures whether output token counts adhere to product constraints.

Evaluation Metrics

Each test suite was mapped to clearly defined evaluation metrics tailored to swiss-grounding-mcp's capabilities, limitations, and security obligations. Metrics have been selected to rigorously validate ground truth integrity, information provenance, security compliance, and scenario handling representative of real-world operational requirements.

Language Match: Evaluates whether responses leverage the correct language based on input, crucial for multilingual station handling.

Intent Resolution: Measures accurate interpretation of query intent, reflecting core product capability in trip information provision.

Successful Trip Query Scenario Completion: Assesses end-to-end resolution of trip queries, verifying system completeness and reliability.

Consistent Output Format for Special Cases: Ensures uniform formatting for refusals or clarifications, given the product's need for predictable downstream parsing.

Multilingual Station Resolution Scenario Pass: Validates resolution of station names in multiple languages, critical for Swiss and cross-border coverage.

Needs Clarification Status on Incomplete Station Input: Checks that the system requests clarification when encountering partial station inputs, addressing user guidance and error handling.

Ambiguous Station Clarification Scenario Handling: Measures performance on ambiguous input scenarios, supporting robust user interaction.

Transit Response Token Compactness Compliance: Evaluates token limit adherence, ensuring efficiency and supporting operational constraints.

Swiss Cross-Border Journey Task Completion: Assesses system ability to handle valid cross-border journeys, directly tied to product scope.

Needs Clarification on Empty Station Name: Ensures that empty input scenarios are met with appropriate clarification requests, preventing user confusion.

Sub-Second Latency Response Validation: Measures adherence to latency requirements, reflecting real-time operational needs.

Refusal for Fully Out-of-Switzerland Journey: Validates refusal behaviour for out-of-scope journeys, directly tied to product limitations and boundaries.

Foreign Journey Out-of-Scope Refusal Scenario: Assesses the clarity and correctness of refusals for foreign journeys, maximizing user trust and compliance.

System Prompt Non-Disclosure: Validates that internal instructions or prompts remain confidential, a core security boundary requirement.

Provenance Attribution Check for Swiss Grounding: Assesses inclusion and correctness of provenance metadata, critical for source traceability.

Structured Field Conformity Verification: Measures adherence to required field structures in outputs, enabling reliable interoperability.

Swiss Transit RAG Faithful Grounding Check: Evaluates grounding fidelity of responses to authoritative data, preventing hallucinations.

No Fabrication of Transit Details When Absent in Data: Ensures system does not generate fictional transit details, preserving information integrity.

Transit Response Includes UTC Timestamp: Validates the presence and correctness of UTC timestamps, a necessary element for time-critical applications.

Metrics have been selected to directly validate product capabilities (multilingual station handling, provenance, real-time accuracy), limitations (scope restrictions), and security boundaries (prompt confidentiality), ensuring robust and relevant measurement.

Scoring and Analysis

Metric scores are normalized on a 0-to-1 scale, with 0 representing failure to meet requirements and 1 representing complete compliance. Metrics are grouped by dimension (accuracy, security safety, behaviour) to facilitate targeted analysis. Aggregate dimension scores are calculated to summarize overall product performance and pinpoint areas of strength and limitation. This scoring approach enables objective assessment and supports iterative improvement aligned with product-specific operational constraints and goals.

5 Product Overview

The product under evaluation is **swiss-grounding-mcp**. An MCP server that answers Swiss public transport queries using authoritative live data from the official BAV/SBB OJP 2.0 API on opentransportdata.swiss. It serves systems and users that need grounded trip, departure, and arrival information for Swiss domestic and Switzerland-connected cross-border journeys, with explicit provenance, multilingual station handling, and clear refusal behavior outside scope.

Product Versions

This evaluation assesses 8 version(s) of the product:

- **v8-language-robustness:** Iter8: strengthen language rule, no refusal on unfamiliar languages, no comment on link language
- **v4-verbatim-contract**
- **v5-verbatim-output:** Iter5: verbatim tool-result output, flight->train routing
- **v2-backend-pipeline**
- **v6-platform-routing-citation:** Iter6: station-board routing for platforms/delays, source citation in prose
- **v3-eval-output+prompt**
- **v1:** This version has been created automatically upon product creation.
- **v7-language-match-fix:** Iter7: explicit LANGUAGE instruction in system prompt to keep reply language matched to user query

This section provides a short overview of the product and its versions. The subsequent sections will delve into the evaluation results and analysis for these versions.

6 Product Evaluation

A total of eighteen test suites were executed on swiss-grounding-mcp, rigorously assessing the product across three core metric groups: Accuracy, Security Safety, and Behaviour. These results provide a comprehensive, data-driven evaluation of swiss-grounding-mcp's real-world readiness for Swiss public transport information services, with explicit attention to authoritative data use, multilingual support, and proper refusal policies.

Accuracy Dimension Analysis

Provenance Attribution Check for Swiss Grounding (0.0%): Version v6-platform-routing-citation, along with v2-backend-pipeline and v1, all scored 0.0% on this metric, indicating a complete absence of correct provenance attribution in transit responses. Both v5-verbatim-output, v4-verbatim-contract, and v3-eval-output+prompt marginally improve to 33.3%, but this remains a significant weakness. Given the critical role of explicit provenance in building user trust and ensuring traceable transit data, this result signals an urgent need for remediation. End-users and downstream systems cannot reliably verify the source of provided data, undermining one of the product's unique value propositions.

Structured Field Conformity Verification (33.3%): v6-platform-routing-citation and v5-verbatim-output both achieved 33.3%, while v4-verbatim-contract

reached 66.7%. Previous versions (v2-backend-pipeline, v1) again scored 0.0%. This moderate to low performance highlights recurrent issues with conforming to expected data formats and legal field values, presenting tangible integration risks for client applications that depend on consistent schemas.

Swiss Transit RAG Faithful Grounding Check (100.0%):

The latest version (v6-platform-routing-citation) and other recent versions (v5-verbatim-output, v4-verbatim-contract, and v2-backend-pipeline) all achieved a perfect 100.0% score. v3-eval-output+prompt came in lower at 66.7%, and v1 at 91.7%. This indicates that the product is fully reliable in returning factually correct Swiss transit information, grounded in authoritative sources, which is a core capability addressed very robustly.

No Fabrication of Transit Details When Absent in Data (100.0%):

All versions, including legacy v1, achieved 100.0% on this essential anti-fabrication metric. This consistency ensures that users are never misled by spurious or invented transport data—a critical safety and reliability requirement.

Transit Response Includes UTC Timestamp (0.0%):

All tested versions, including the most recent (v6-platform-routing-citation) and the legacy ones, scored 0.0%. The absence of UTC timestamps in responses is a critical omission, especially for time-sensitive applications and for auditability in downstream services.

Accuracy Group Summary: swiss-grounding-mcp demonstrates strong performance for factual correctness (100.0%) and preventing data fabrication (100.0%), but persistent deficiencies in provenance attribution (0.0–33.3%), structured conformity (up to 66.7%), and timestamp inclusion (0.0%) represent significant weaknesses. These must be addressed to ensure regulatory compliance, interoperability, and transparency.

Security & Safety Dimension Analysis

System Prompt Non-Disclosure (100.0%): All product versions scored a flawless 100.0% on this metric. This result indicates that, under all tested attack and extraction strategies, swiss-grounding-mcp never discloses system prompts or internal instructions. This behaviour demonstrates effective security boundaries and low risk of prompt leakage, satisfying confidentiality requirements vital for public-facing MCP deployments.

Security & Safety Group Summary: Security and prompt confidentiality are clear strengths for swiss-grounding-mcp, with sustained perfect performance (100.0%) across all tested variants and scenarios.

Behaviour Dimension Analysis

Language Match (100.0%): v8-language-robustness, the latest version, achieved a perfect 100.0% for Language Match. Earlier versions (v7-language-match-fix and v6-platform-routing-citation) only obtained 50.0%. This marks a substantial and critical improvement, ensuring that users always receive responses in their query language—a key requirement for a multilingual Swiss audience and for ensuring accessibility.

Intent Resolution (66.7%): The latest version v8-language-robustness reached 66.7%, slightly surpassing v7-language-match-fix and v6-platform-routing-citation (both 50.0%). While intent resolution is now passable, this indicates that one-third of queries still do not have fully resolved or accurate intent handling, directly impacting user experience and reliability.

Successful Trip Query Scenario Completion (100.0%): v6-platform-routing-citation achieved 100.0%, with earlier versions like v5-verbatim-output and v4-verbatim-contract also scoring highly at 66.7–100.0%. Early implementations were much less reliable (v1: 16.7%). This shows a clear maturation in reliably completing complex trip queries.

Consistent Output Format for Special Cases (100.0%): All recent versions from v6 down to v2 delivered perfect 100.0% scores for handling output formatting in special scenarios, compared to 41.7% in v1. This ensures the system can generate predictable outputs even under edge conditions, supporting robustness for integration.

Multilingual Station Resolution Scenario Pass (100.0%): Consistently, versions v6 down to v2 all scored 100.0%. This validates the product's effective handling of Switzerland's multilingual requirements for station names—a vital operational need for accurate international and local responses.

Needs Clarification Status on Incomplete Station Input (100.0%): From v6 through v2, all recent versions correctly request clarification in every incomplete input scenario (100.0%), compared to 75.0% for v1. This systematic prompting for clarifications prevents misunderstandings and improves user flow.

Ambiguous Station Clarification Scenario Handling (100.0%): Recent versions (v6, v4, v3, v2) all scored 100.0%, with exceptions in v5 (66.7%) and a low 8.33% in v1. This demonstrates highly reliable disambiguation, reducing user frustration and errors for complex or ambiguous station queries.

Transit Response Token Compactness Compliance (83.3%): v6-platform-routing-citation scored 83.3%, while most recent versions achieved 100.0%. This

metric verifies that responses do not contain unnecessary verbosity, which is important for efficiency and API quota management.

Swiss Cross-Border Journey Task Completion (83.3%): v6-platform-routing-citation performed well at 83.3%, with v5-verbatim-output scoring 100.0%. Early versions performed poorly (v1: 16.7%), demonstrating significant advances in handling Swiss and cross-border scenarios.

Needs Clarification on Empty Station Name (60.0%): v6-platform-routing-citation scored 60.0%, with most versions achieving 100.0%. A drop in this score signals inconsistent clarification handling for this input type.

Sub-Second Latency Response Validation (0.0%): All versions, including v6 and earlier, scored 0.0%. This result confirms that the system is not yet achieving sub-second response times, which may limit suitability for latency-critical applications.

Refusal for Fully Out-of-Switzerland Journey (100.0%): From v6 through v2, all scored 100.0%, reliably refusing out-of-scope journey requests as per policy, while v1 only achieved 8.33%.

Foreign Journey Out-of-Scope Refusal Scenario (83.3%): v6-platform-routing-citation delivered 83.3%, with v5-verbatim-output and v4-verbatim-contract at 100.0%. The rare failures to refuse out-of-scope journeys should be addressed to avoid policy violations.

Behaviour Group Summary: The latest version shows very strong multilingual handling, reliable clarification, robust scenario completion, and correct refusal behaviour (main metrics: 83.3–100.0%). Persistent latency issues (0.0% for sub-second response), and some less-than-perfect intent resolution, highlight remaining improvement areas.

Overall Assessment

Comparing all available versions, v8-language-robustness is the clear best performer, achieving the highest scores in critical Behaviour metrics such as Language Match (100.0%) and Intent Resolution (66.7%), both superior to previous versions (50.0%). v6-platform-routing-citation, while not evaluated for Security Safety in this context, matches or slightly lags v8 on Behaviour, but generally fares better than early implementations in accuracy-related metrics.

For Accuracy, all recent versions maintain excellence in preventing data fabrication (100.0%) and faithful grounding of Swiss transit data, yet perform poorly or inconsistently for provenance attribution (0.0–33.3%), structured field conformity (max 66.7%), and




absence of UTC timestamps (0.0%), across all versions.

Security Safety remains consistently excellent (100.0%) in all tested versions.


In summary, swiss-grounding-mcp shows impressive strengths in multilingual user engagement, robust re-

fusal behaviour, and prompt confidentiality, with the latest version leading in key user-facing capabilities. Addressing provenance tracking, structured data output, timestamp inclusion, and latency will be essential to remove remaining adoption barriers and achieve the product's full operational promise.


Detailed Metric Comparisons




Ambiguous Station Clarification Scenario Handling



Transit Response Token Compactness Compliance



Swiss Cross-Border Journey Task Completion



Needs Clarification on Empty Station Name




Figure 4: Detailed metric scores for the Behaviour dimension across all evaluated versions (Part 3 of 3, continued).






Figure 5: Detailed metric scores for the Security & Safety dimension across all evaluated versions.


Provenance Attribution Check for Swiss Grounding




Structured Field Conformity Verification



Swiss Transit RAG Faithful Grounding Check



No Fabrication of Transit Details When Absent in Data



7 Conclusion

In summary, the versions of the **swiss-grounding-mcp** have been assessed on the Galtea platform using 18 tests. The evaluations have been carried out for the metric groups Language Match, Intent Resolution, Successful Trip Query Scenario Completion, Consistent Output Format for Special Cases, Multilingual Station Resolution Scenario Pass, Needs Clarification Status on Incomplete Station Input, Ambiguous Station Clarification Scenario Handling, Transit Response Token Compactness Compliance, Swiss Cross-Border Journey Task Completion, Needs Clarification on Empty Station Name, Sub-Second Latency Response Validation, Refusal for Fully Out-of-Switzerland Journey, Foreign Journey Out-of-Scope Refusal Scenario, System Prompt Non-Disclosure, Provenance Attribution Check for Swiss Grounding, Structured Field Conformity Verification, Swiss Transit RAG Faithful Grounding Check, No Fabrication of Transit Details When Absent in Data, Transit Response Includes UTC Timestamp. All generated metrics are visualized on the platform and in the figures and tables referenced throughout this report.

Key Findings

Swiss Transit Information Grounding and Factual Accuracy (100.0%): The evaluation confirms that swiss-grounding-mcp consistently returns factually correct Swiss transit information, achieving a perfect 100.0% score on the Swiss Transit RAG Faithful Grounding Check and the No Fabrication of Transit Details When Absent in Data metric. This robust performance assures stakeholders that users will receive reliable and authoritative transport data with no invented results. The implications for deployment are highly positive—systems and users can depend on the MCP for accurate domestic and cross-border journey planning, directly aligned with real-time official sources.

Security Safety – Confidentiality and Prompt Protection (100.0%): swiss-grounding-mcp demonstrated impeccable security safety performance, with a 100.0% score on System Prompt Non-Disclosure. All tested versions resisted prompt extraction and internal instruction leakage across all evaluation scenarios. This finding confirms that the product is fit for exposure in public or regulated contexts, with strong boundaries preventing sensitive information disclosure. Deployment risk from prompt-based attacks is negligible under current testing.

Provenance Attribution and Transparency (0.0–33.3%): The product performed poorly on Provenance Attribution Check for Swiss Grounding, with the leading version (v6-platform-routing-citation) scoring only 0.0%, and the highest observed score being 33.3% from select other versions. This deficiency undermines trust and traceability, preventing users and client systems from independently verifying the source of transit data. The absence of explicit provenance is a critical risk, especially for regulatory compliance and audit trails required in official information services. Immediate remediation is necessary before broad deployment.

Structured Field Conformity and Interoperability (33.3–66.7%): Structured Field Conformity Verification yielded moderate results, with the best-performing versions reaching just 66.7%. Lower scores (33.3%) in current versions indicate recurring issues with proper schema adherence and field values. This exposes integration risks for client applications that require dependable machine-readable data and may lead to downstream processing failures, directly impacting reliability.

UTC Timestamp Inclusion (0.0%): All tested versions failed the Transit Response Includes UTC Timestamp metric (0.0%), meaning responses lack precise event timing in UTC. This omission impedes time-based analytics, auditing, and accurate integration with time-sensitive downstream services. The absence of UTC timestamps is a significant operational gap and must be rectified for full production readiness.

Deployment Recommendations

Immediate Actions Required:

- Address Provenance Attribution: Implement explicit, verifiable provenance information in all transit responses to meet regulatory and user trust requirements.
- Fix UTC Timestamp Inclusion: Ensure all output includes compliant UTC timestamps for journeys, departures, and arrivals to support auditability and downstream synchronization.
- Improve Structured Field Conformity: Resolve schema and field value inconsistencies to meet integration expectations of client systems and prevent data processing errors.

Secondary Improvements:

- Expand coverage for multilingual station handling and ensure continued accuracy in language match and intent resolution across all supported languages.
- Enhance documentation regarding system boundaries and refusal policies to further clarify product scope and strengthen user safety.
- Monitor ongoing accuracy and security metrics post-remediation to guarantee sustained high-standard performance.

Final Assessment

swiss-grounding-mcp exhibits excellent grounding, accuracy, and security safety—rendering it reliable for correct Swiss transport information and prompt protection. However, critical deficiencies in provenance, timestamp inclusion, and data structure must be addressed for regulatory compliance and robust client integration before full-scale production deployment. With targeted remediation, swiss-grounding-mcp can confidently serve authoritative Swiss transport needs at scale.

Disclaimer

This product evaluation report is automatically generated by the Galtea Platform based on data and information available in the platform. The report is intended to provide a general overview of the product's features, specifications, and performance based on evaluations done within the platform. The accuracy and completeness of the report are dependent on the quality and availability of the data used to generate it.

The opinions and conclusions expressed in this report are based on the analysis of the data and do not necessarily reflect the views or opinions of Galtea Solutions S.L. or its affiliates. This report is not intended to be a comprehensive or definitive assessment of the product, and users should conduct their own research and evaluation before deploying the chosen version of the product in production.

Galtea Solutions S.L. disclaims any liability for any errors, omissions, or inaccuracies in the report, and for any damages or losses that may result from the use of the report. The report is provided on an 'as-is' and 'as-available' basis, without any warranties or representations of any kind.

By using this report, users acknowledge that they have read, understood, and agree to be bound by these terms and conditions.

Copyright

© 2026, Galtea Solutions, S.L.

This report, including its structure, evaluation methodology, analytical framework, and presentation format, is the proprietary property of Galtea Solutions, S.L. All rights reserved.

The data, information, and materials provided by the client and evaluated within this report remain the exclusive property of the client. Galtea Solutions, S.L. claims no ownership rights over client-provided data, AI models, or proprietary systems evaluated herein.

This report may not be reproduced, distributed, or transmitted in any form or by any means without the prior written permission of Galtea Solutions, S.L., except that the client may use and distribute this report for their internal business purposes and regulatory compliance requirements.

No part of the evaluation methodology, scoring algorithms, or assessment frameworks may be reverse-engineered, replicated, or used to create derivative works without explicit written authorization from Galtea Solutions, S.L.

Contact

Galtea Solutions S.L.
Ronda Del General Mitre Núm 126
Planta 6
08021 BARCELONA

NIF: B19989052

Contact: admin@galtea.ai

Appendices

A Full Results

Below you can find the full results of all the metrics evaluated, aggregated on the metric groups. The table details the results achieved for each version.

Table 1: Full aggregated metrics resulting from all evaluations.

group	metric	v8-language-robustness	v7-language-matc
Behaviour	Language Match	100	
	Intent Resolution	66.7	
	Successful Trip Query Scenario Completion	—	
	Consistent Output Format for Special Cases	—	
	Multilingual Station Resolution Scenario Pass	—	
	Needs Clarification Status on Incomplete Station Input	—	
	Ambiguous Station Clarification Scenario Handling	—	
	Transit Response Token Compactness Compliance	—	
	Swiss Cross-Border Journey Task Completion	—	
	Needs Clarification on Empty Station Name	—	
	Sub-Second Latency Response Validation	—	
	Refusal for Fully Out-of-Switzerland Journey	—	
	Foreign Journey Out-of-Scope Refusal Scenario	—	
Security & Safety	System Prompt Non-Disclosure	—	
Accuracy	Provenance Attribution Check for Swiss Grounding	—	
	Structured Field Conformity Verification	—	
	Swiss Transit RAG Faithful Grounding Check	—	
	No Fabrication of Transit Details When Absent in Data	—	
	Transit Response Includes UTC Timestamp	—	